

Deep Learning for NLP

Universität Bielefeld

Lecture 4 – Word Embeddings 1: Word2Vec


Dr. Steffen Eger

steffen.eger@uni-bielefeld.de




Previous lectures

- How do we build/train neural networks?
- Today and next two weeks: How do we represent the input?
- Representing images:



<http://hosting.soonet.ca/eliris/remotesensing/LectureImages/pixel.gif>



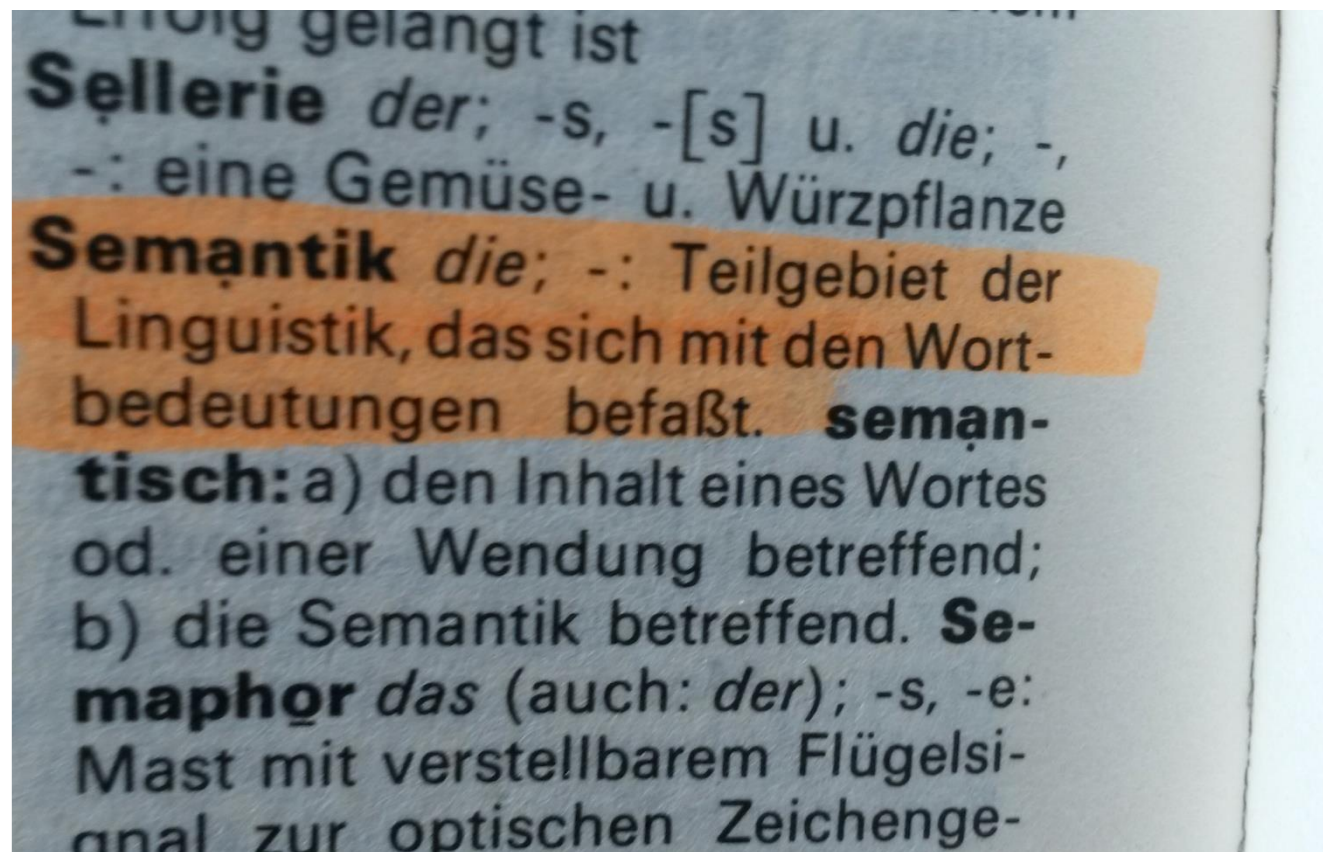
http://www.comp.leeds.ac.uk/vision/vision_matlab_tutorial/images/section2/MatRGB.jpg

- How can we represent words?

Word meaning


How can we represent words?

- As a dictionary entry



Taxonomy of words

- Represent words by their **relations to other words**



Picture from: <http://kylescholz.com/projects/wordnet/>, based on representation from WordNet: <https://wordnet.princeton.edu>

- “One-hot” vector, sparse representation

der	die	und	in	für
1	0	0	0		0
0	1	0	0		0
0	0	1	0		0
0	0	0	1		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		1
...

- Dimensionality of vector equals size of vocabulary

- “One-hot” vector, sparse representation

der	die	und	in	...	für
1	0	0	0		0
0	1	0	0		0
0	0	1	0		0
0	0	0	1		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		0
0	0	0	0		1
...

Problem: relations between words are not represented

Word meaning can be inferred from context

- Famous example by McDonald and Ramscar (2001):
 1. He filled the **wampimuk**, passed it around and we all drank some.

Word meaning can be inferred from context

- Famous example by McDonald and Ramscar (2001):

1. He filled the **wampimuk**, passed it around and we all drank some.

jar

cup

glass

goblet

...

Word meaning can be inferred from context

- Famous example by McDonald and Ramscar (2001):

1. He filled the wampimuk, passed it around and we all drank some.

jar

cup

glass

goblet

...

2. We found a little hairy **wampimuk** sleeping behind the tree.

Word meaning can be inferred from context

- Famous example by McDonald and Ramscar (2001):

1. He filled the wampimuk, passed it around and we all drank some.

jar

cup

glass

goblet

...

2. We found a little hairy **wampimuk** sleeping behind the tree.

cat

bear

raccoon

mole

...

- Firth (1957): *“You shall know a word by the company it keeps.”*

Computational semantics: Count models

How can we model the distributional hypothesis?

- By calculating co-occurrence counts
 - capture in which contexts a word appears
- Context is modeled using a window over the words
- Consider the following example (from Richard Socher's lecture):
- Corpus
 - *I like deep learning .*
 - *I like NLP .*
 - *I enjoy flying .*
- Window size = 1, left and right neighbor
 - In real tasks, window size is usually bigger (5-10)

How can we model the distributional hypothesis?

- By calculating co-occurrence counts
 - capture in which contexts a word appears
- Context is modeled using a window over the words
- Consider the following example by Rich
- Corpus
 - *I like deep learning .*
 - *I like NLP .*
 - *I enjoy flying .*
- Window size = 1, left and right neighbor
 - In real tasks, window size is usually bigger (5-10)

Such models have been called „count models“ in the literature

See: Baroni et al. Don't count predict! A systematic comparison of context-counting vs. context-predicting semantic vectors. In: ACL 2014

Co-occurrence matrix

- Example by Richard Socher:
 - *I like deep learning .*
 - *I like NLP .*
 - *I enjoy flying .*

Window size = 1

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0

Co-occurrence matrix

- Example by Richard Socher:

- I like deep learning .*
- I like NLP .*
- I enjoy flying .*

Window size = 1

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0

Co-occurrence matrix

- Example by Richard Socher:

- I like deep learning .*
- I like NLP .*
- I enjoy flying .*

Window size = 1

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0

Co-occurrence matrix

- Example by Richard Socher:
 - *I like deep learning .*
 - *I like NLP .*
 - *I enjoy flying .*

Window size = 1

counts	I	like	enjoy	deep	learning	NLP	flying	.
I	0	2	1	0	0	0	0	0
like	2	0	0	1	0	1	0	0
enjoy	1	0	0	0	0	0	1	0
deep	0	1	0	0	1	0	0	0
learning	0	0	0	1	0	0	0	1
NLP	0	1	0	0	0	0	0	1
flying	0	0	1	0	0	0	0	1
.	0	0	0	0	1	1	1	0

Co-occurrence counts

- Assumption: If we collect co-occurrence counts over thousands of sentences, the vectors for “enjoy” and “like” will have similar vector representations.

Co-occurrence counts

- Assumption: If we collect co-occurrence counts over thousands of sentences, the vectors for “enjoy” and “like” will have similar vector representations.
- Problem:
 - Vectors become very large with real data
→ We need to apply dimensionality reduction

Computational semantics: NN models

Background idea: language models

- Based on the concept of **language modeling**:
- Common problem in NLP, popular application is auto-completion
 - Given a sequence of words, predict the following word
 - *The same procedure as every _____*
- Idea:
(Classical) Language modeling is too restrictive because it only considers the left context. What about the right context?

- Most popular toolkit for training word representations:

word2vec


<https://code.google.com/archive/p/word2vec/>

Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado, and Jeffrey Dean:
Distributed Representations of Words and Phrases and their Compositionality
In Proceedings of NIPS, 2013.


- Two different (language modeling) auxiliary tasks: CBOW and Skip-gram

- CBOW: Given a context, predict the missing word
 - *same procedure ____ every year*
 - *as long ____ you sing*
 - *please stay ____ you are*
- Skip-gram: given a word, predict the context words
 - ____ *as* ____
 - If window size is two, we aim to predict:
 (w, c_{-2}) , (w, c_{-1}) , (w, c_1) and (w, c_2)

CBOW vs Skip-gram




CBOW

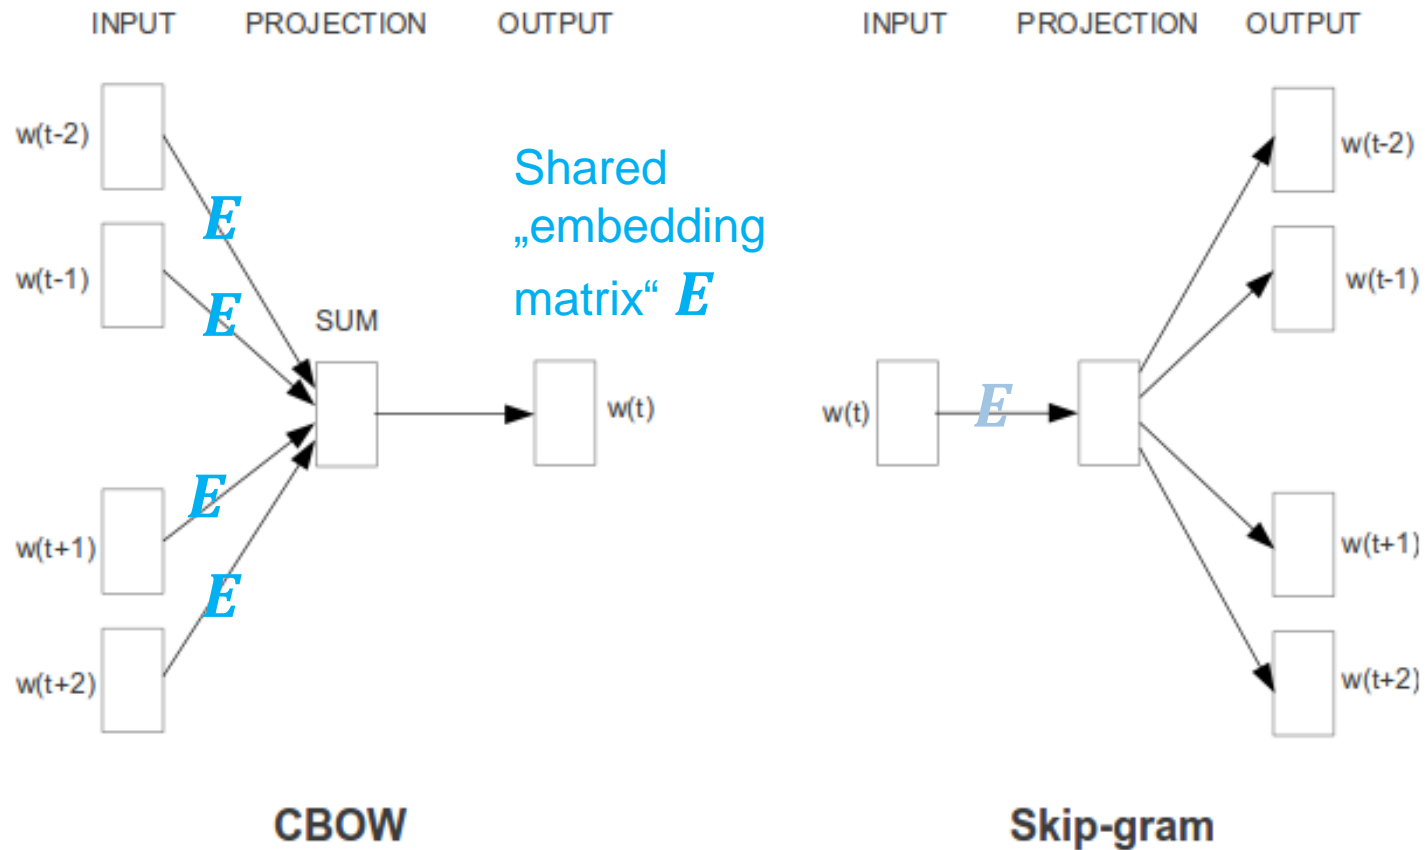


Skip-gram

CBOW vs Skip-gram



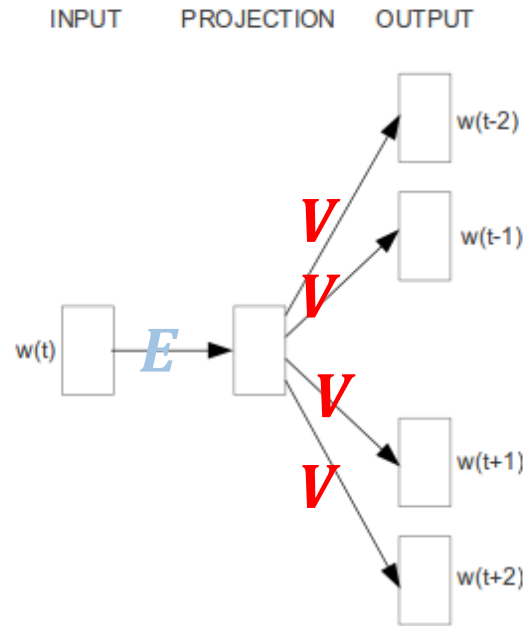
CBOW vs Skip-gram



The Skip-gram model: Theory

We'll take a closer look at the Skip-Gram model

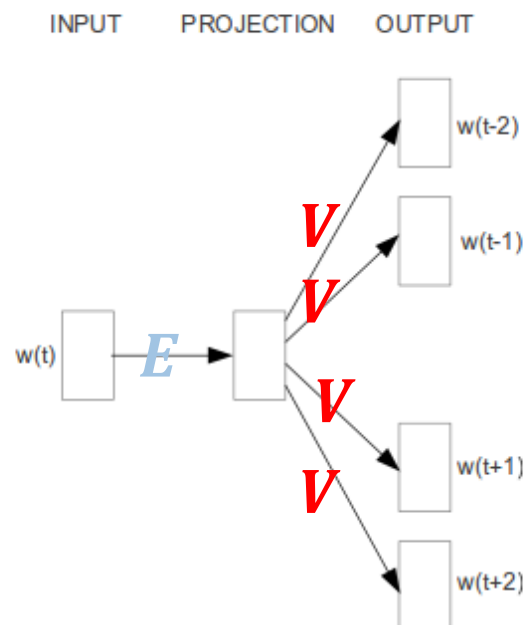
- E has dimension $N \times d$
- N is number of words in the vocabulary
- d is the embedding dimension
- V has dimension $d \times N$



Skip-gram

The Skip-gram model: Theory

- We'll take a closer look at the Skip-Gram model




How can we predict different words when V is always the same?

Skip-gram

The Skip-gram model: Theory

We'll take a closer look at the Skip-Gram model



How can we predict different words when V is always the same?

Turns out the original model is actually this:

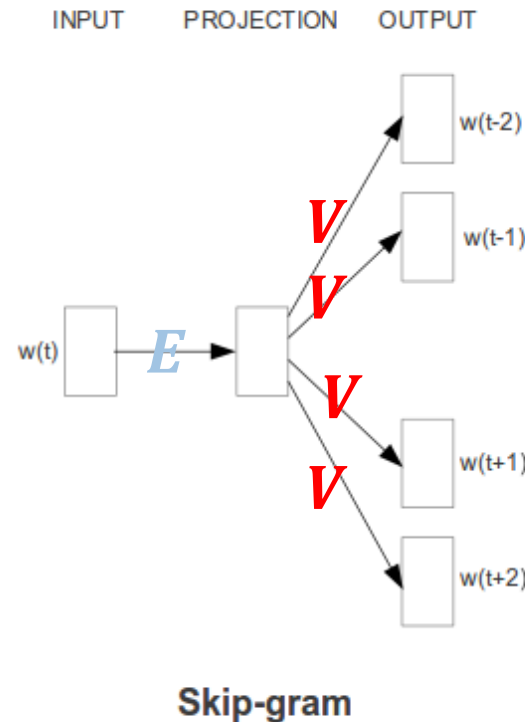
Skip-gram



The Skip-gram model: Theory

We'll take a closer look at the Skip-Gram model

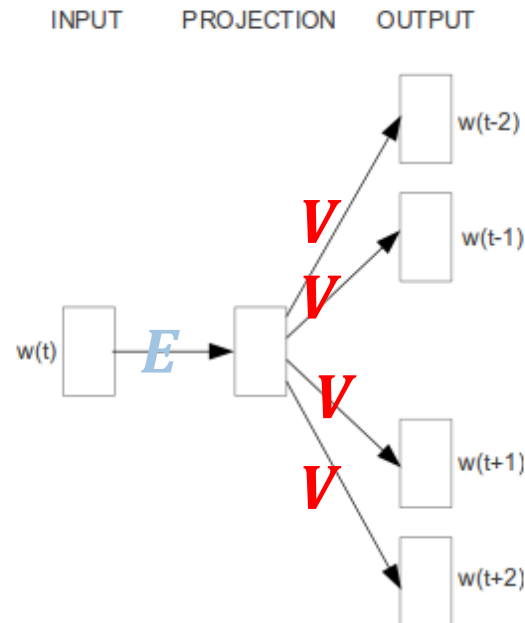
- $e = e(w) = wE$ has dimension $1 \times d$: It's the *embedding* of word w
- The activation of the projection layer is *linear* (= *identity*: $f(x)=x$)
- eV has dimension $1 \times N$
- $V = [v_1 \cdots v_N]$: each v_i can be seen as an(other) embedding of a vocab. word



The Skip-gram model: Theory

We'll take a closer look at the Skip-Gram model

- $eV = [ev_1, \dots, ev_N]$ has dimension $1 \times N$
- $V = [v_1, \dots, v_N]$: each v_i can be seen as an(other) embedding of a vocab. word
- The output layer has softmax activation function
- $\text{softmax}(eV) = [\exp(ev_1), \dots, \exp(ev_N)] / Z$, where Z is normalizer




Skip-gram

The Skip-gram model: Theory

We'll take a closer look at the Skip-Gram model

- Could just run this model with SGD
- With methods we learned
- After training, we're interested in the E matrix, which holds the word embeddings
- **However, the practical implementation is different from this (see below)**



Skip-gram

Discuss two (or more) limitations of the Word2Vec model



The Skip-gram model: Illustration

- Preparation:
 - Download a lot of (**unlabeled**) data, e.g. all the poems of W.S.

*All the world's a stage and all the men and women merely players.
They totter ...*

- Tokenize it

*All the world 's a stage and all the men and women merely players .
They totter ...*

- For word2vec: Define either one of two **auxiliary** tasks: predict middle words or predict contexts

For skip-gram:

*All the world 's a stage and all the men and women merely players .
They totter ...*

- Training data (maybe this is our first batch):
 - x=world t=the
 - x=world t=All
 - x=world t='s
 - x=world t=a
- Of course, the 1-hot vectors of this

For skip-gram:

*All the world's a stage and all the men and women merely players .
They totter ...*

- Training data (maybe this is our first batch):

- $x=\text{world}$ $t=\text{the}$
- $x=\text{world}$ $t=\text{All}$
- $x=\text{world}$ $t=\text{'s}$
- $x=\text{world}$ $t=\text{a}$

Feed in to the network; update params



- Of course, the 1-hot vectors of this

For skip-gram:

*All the world's a stage and all the men and women merely players .
They totter ...*

- Training data (maybe this is our first batch):

- $x=\text{world}$ $t=\text{the}$
- $x=\text{world}$ $t=\text{All}$
- $x=\text{world}$ $t=\text{'s}$
- $x=\text{world}$ $t=\text{a}$

Feed in to the network; update params



- Of course, the 1-hot vectors of this
- Notation: We write (x,t) or (w,c) for a training data sample

For skip-gram:

*All **the world's** **a stage** and all the men and women merely players .
They totter ...*

- Training data (maybe this is our second batch):

- $x='s$ $t=world$
- $x='s$ $t=the$
- $x='s$ $t=a$
- $x='s$ $t=stage$

Feed in to the network; update params



- Of course, the 1-hot vectors of this
- Notation: We write (x,t) or (w,c) for a training data sample

For skip-gram:

*All the world's a stage and all the men and women merely players .
They totter ...*

Context window size is a hyperparam.

- Training data (maybe this is our second batch):

- $x='s$ $t=world$
- $x='s$ $t=a$

Feed in to the network; update params



- Of course, the 1-hot vectors of this
- Notation: We write (x,t) or (w,c) for a training data sample

For skip-gram:

*All the world 's a stage and all the men and women merely players .
They totter ...*

Context window size is a hyperparam.

- Training data (maybe this is our second batch):

- $x='s$ $t=All$
- $x='s$ $t=the$
- $x='s$ $t=world$
- $x='s$ $t=a$
- $x='s, t=stage, x='s, t=and$

Feed in to the network; update params



- Of course, the 1-hot vectors of this
- Notation: We write (x,t) or (w,c) for a training data sample

What if the window size is arbitrarily large in Skip-Gram?

A: this can only work when embedding size is huge

B: quality of representations decreases

C: quality of representations increases

D: all words will get same embeddings



Skip-gram: practical realization


Problem:

- Computing the softmax in the output layer is very costly, particularly when the size of the vocabulary is large
- Because we have to normalize

Skip-gram: practical realization

Idea: replace softmax by sigmoid activation function in the output layer


- No normalization \rightarrow cheaper
- So, at each step in the optimization:
 - input vector is a 1-hot vector of the center word w
 - Truth is a 1-hot vector of target/context word c
 - As loss, we choose cross-entropy loss: $-\sum_i t_i \log y_i$
 - $y_i = \sigma(e(w) \cdot v_i)$



- **But now a different problem pops up:**
 - What if our model learns that E and V are matrices with huge entries?
 - Then $y_i = 1$ for all output units and our loss becomes 0, which is the optimum
 - Note that this would not happen with softmax, because $y_i = 1/N$ in this case
- **Solution:**
 - We choose random words (*negative samples*)
 - For an input (w, c) , we set the target vector as a vector with 1's at the position of c and the positions of the random words
 - Choose the activations of the negative samples as $f(x) = 1 - \sigma(x) = \sigma(-x)$
 - When $f(x) = 1$, then $\sigma(x) = 0$

Skip-gram: practical realization

- Can think of it as follows:



$$\text{Loss} = - \sum_i t_i \log y_i$$

- With cross-entropy loss, the loss becomes:

$$- [\log(\sigma(\mathbf{e}(\mathbf{w}) \cdot \mathbf{v}_c)) + \sum_r \log(\sigma(-\mathbf{e}(\mathbf{w}) \cdot \mathbf{v}_r))]$$

$$\sigma(-x) = 1 - \sigma(x)$$

- r are the negative samples
- This is not quite a standard neural network, but close

Toolkits for training word representations

word2vec

<https://code.google.com/archive/p/word2vec/>

GloVe

<http://nlp.stanford.edu/projects/glove/>

- GloVe aims at reconciling the advantages of global co-occurrence counts and local context windows
- Applies additional trick: take the sum of the target/center vector $\mathbf{e}(\mathbf{w})$ and the context vector \mathbf{v}_c of each word as representation
- Many more, but these are two popular ones
- Terminology:
 - **word representations \approx word embeddings \approx word vectors**
 - context-counting vs context-predicting representations, sparse vs dense

Pre-Trained Embeddings

- Word2vec
 - trained on Google news (100 billion tokens)
 - vectors with Freebase naming, trained on news (100 billion tokens)
- GloVe
 - trained on Wikipedia (6 billion tokens)
 - trained on CommonCrawl (42 and 840 billion tokens)
 - trained on Twitter (27 billion tokens)
- Omer Levy: dependency-based embeddings trained on Wikipedia
<https://levyomer.wordpress.com/2014/04/25/dependency-based-word-embeddings/>
- There are many embeddings nowadays, in all possible languages
<https://fasttext.cc/docs/en/crawl-vectors.html>

Evaluation of word embeddings

The look and feel of word representations

...

wiegen 0.0427915 -0.401344 -0.0667862 0.21649 0.00169907 -0.0687786 -0.195405 -
0.0534437 -0.676733 -0.23975 -0.159674 -0.0402676 -0.0617923 0.284718 0.358523 -
0.285709 -0.00736682 -0.254635 -0.22907 -0.186109 ...

einlegen 0.365857 -0.0339146 0.198442 -0.0961315 0.156193 0.253468 0.169963 -
0.232588 -0.422901 -0.0750184 0.0236783 0.249385 -0.0122247 -0.584567 -0.0711365
0.254896 0.382103 0.352294 0.825432 0.277691 0.773015 ...

sprengstoff 0.06961 0.118456 0.00497905 0.581913 -0.326157 -0.0674812 -0.0926074 -
0.254514 -0.458406 -0.225093 0.0424881 -0.142328 -0.138707 0.481305 0.183707 -
0.626077 0.396159 0.156636 0.157851 -0.441935...

vulkan 0.0322022 -0.429981 0.352328 -0.0530384 -0.366048 0.44187 -0.265227 -0.223954
-0.369078 -0.203064 0.158458 0.169517 0.448234 0.497058 -0.20855 0.046978 0.180444
0.290595 0.00907329 0.130582 0.0378717 -0.339296 0.399039...

...

- Extrinsic

by the performance of a model that uses the word representations for solving a task

- Named entity recognition (accuracy), machine translation (BLEU score), summarization (ROUGE score), information retrieval (coverage)...
- Compare performance of two models that only differ in the word representations they use

- Intrinsic

by using the representations directly

- Word Similarity Task
- Word Analogy Task
- Word Intrusion Task

Extrinsic Evaluation: Setup

- Say, our task is POS tagging
- Our labeled training data

Word	Label
The	DET
cat	NN
on	PREP
the	DET
mat	NN
.	PUNC

Extrinsic Evaluation: Setup

- Say, our task is POS tagging
- Our labeled training data; **replace words with their embeddings**

x	t
E(The)	1-hot(DET)
E(cat)	1-hot(NN)
E(on)	1-hot(PREP)
E(the)	1-hot(DET)
E(mat)	1-hot(NN)
E(.)	1-hot(PUNC)

Extrinsic Evaluation: Setup

- Say, our task is POS tagging
- Our labeled training data; **replace words with their embeddings; usually add some context**

x	t
E(SOS);E(The);E(cat)	1-hot(DET)
E(The);E(cat);E(on)	1-hot(NN)
E(cat); E(on); E(on)	1-hot(PREP)
E(on); E(the); E(mat)	1-hot(DET)
E(the); E(mat); E(.)	1-hot(NN)
E(mat); E(.); E(EOS)	1-hot(PUNC)

Assume you would use an MLP as a model in the previous setting. How would inputs, outputs, etc. look like?



Extrinsic Evaluation: Setup

- Say, our task is POS tagging
- Our labeled training data; **replace words with their embeddings; usually add some context** [not necessary with other architectures such as RNN; see later lectures]

x	t
E(SOS);E(The);E(cat)	1-hot(DET)
E(The);E(cat);E(on)	1-hot(NN)
E(cat); E(on); E(on)	1-hot(PREP)
E(on); E(the); E(mat)	1-hot(DET)
E(the); E(mat); E(.)	1-hot(NN)
E(mat); E(.); E(EOS)	1-hot(PUNC)


- Now train with different embeddings and look which one performs best

- Similar words should have similar representations
 - Similarity Dataset: <http://www.cs.technion.ac.il/~gabr/resources/data/wordsim353/>
Scores from 0 to 10 by human raters
 - Intrinsic evaluation of embeddings:
 - quantify similarity by similarity of word vectors
 - evaluate correlation with human judgements

Word 1	Word 2	Human (mean)	Learned vectors
tiger	cat	7.35	cossim(tiger, cat)
book	paper	7.46	cossim(book, paper)
plane	car	5.77	cossim(plane, car)
smart	student	4.62	cossim(smart, student)
stock	phone	1.62	cossim(stock, phone)
....		

Relations between word vectors

- Mikolov et al. (2013)



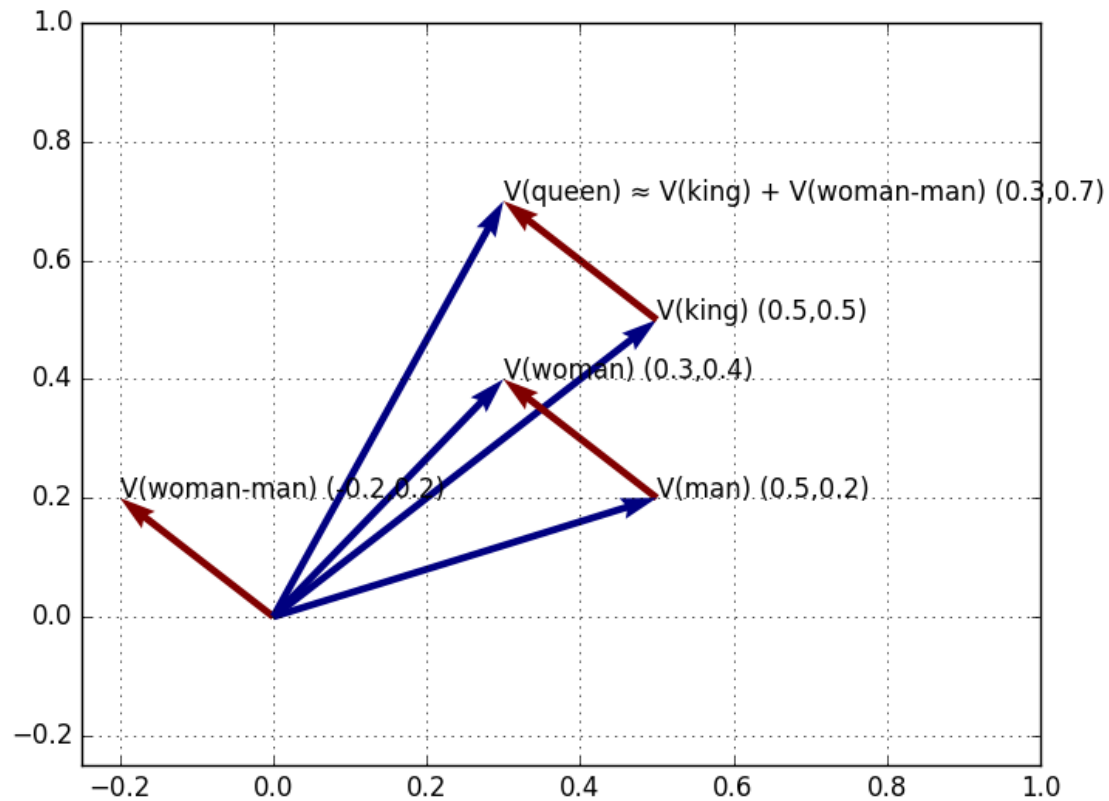
How to find analogies?

- A is to B as C to ?
 - Germany is to Berlin as France to x
- Find x such that:
 - $\text{vec}(x) = \text{vec}(\text{"Berlin"}) - \text{vec}(\text{"Germany"}) + \text{vec}(\text{"France"})$

- Most famous example:

KING – MAN + WOMAN = QUEEN

KING-MAN+WOMAN=QUEEN



Semantic analogies

- All examples from:
<https://code.google.com/p/word2vec/source/browse/trunk/questions-words.txt>
- capital-common-countries
 - *Athens Greece Baghdad* ***Iraq***
 - *Athens Greece Berlin* ***Germany***
- currency
 - *Denmark krone Croatia* ***kuna***
 - *Europe euro Hungary* ***forint***
- family
 - *boy girl brother* ***sister***
 - *brother sister dad* ***mom***

Syntactic analogies

- adjective-to-adverb
 - *amazing amazingly apparent* ***apparently***
- comparative
 - *bad worse big* ***bigger***
- present-participle
 - *code coding dance* ***dancing***
- past-tense
 - *dancing danced decreasing* ***decreased***
- plural
 - *banana bananas bird* ***birds***
- 3rd person verbs
 - *decrease decreases eat* ***eats***

Try it out!

- word2vec code: ./demo-analogy.sh

```
Enter three words (EXIT to break): loud louder slow
Word: loud   Position in vocabulary: 9481
Word: louder Position in vocabulary: 18502
Word: slow   Position in vocabulary: 2188

-----
Word                Distance
-----
faster              0.546676
slower              0.545372
efficient           0.445998
downside            0.426637
cheaper             0.419650
slowly              0.418785
slowing             0.418599
gradual             0.417532
quicker             0.404312
```


- Word2Vec and Glove are pretty good tools
- Fast, give good word embeddings
- However, many other embeddings out there (see next lectures)
- Always try out different embeddings --- consider them as another hyperparameter
 - Results may vary drastically with different embeddings

When and why are embeddings helpful?



- Vectors are useful for representing words
 - Dense vs sparse representations
 - Projecting co-occurrence counts to low-dimensional vectors vs directly learning low-dimensional vectors
- Learning low-dimensional vectors
 - Inspired by neural language modeling
 - CBOW and Skip-gram model
 - Negative sampling
- Evaluating word representations
 - Extrinsic vs intrinsic evaluation
- Terminology:
word representations \approx word embeddings \approx word vectors

References

- Levy & Goldberg: word2vec Explained: deriving Mikolov et al.'s negative-sampling word-embedding method, preprint on arxiv: *arXiv:1402.3722*, 2014.
- Mikolov, T., Sutskever, I., Chen, K., Corrado, G. S., & Dean, J.: Distributed representations of words and phrases and their compositionality, in *Advances in neural information processing systems*: 3111-3119, 2013.
- McDonald, S., & Ramscar, M.: Testing the distributional hypothesis: The influence of context on judgements of semantic similarity, in *Proceedings of the 23rd annual conference of the Cognitive Science Society*, 2001.
- Bengio, Y., Schwenk, H., Senécal, J. S., Morin, F., & Gauvain, J. L.: Neural probabilistic language models, in *Innovations in Machine Learning*: 137-186, 2006.
- Collobert, R., & Weston, J.: A unified architecture for natural language processing: Deep neural networks with multitask learning, in *Proceedings of the 25th international conference on Machine learning*: 160-167, 2008.
- Derweester, S., Dumais, S.T., Furnas, G.W., Landauer, T.K., & Harshman, R.: Indexing by latent semantic analysis , in *Journal of the American Society for Information Science* 41: 391-407, 1990.
- Harris, Z. S.: Distributional structure. *Word* 10(2-3): 146-162, 1954.
- John R. Firth: A synopsis of linguistic theory 1930–55, in *Studies in Linguistic Analysis* (special volume of the Philological Society), 1–32, 1957.

- Baroni and, M., Dinu, B. & Kruszewski, G.: Don't count, predict! A systematic comparison of context-counting vs. context-predicting semantic vectors, in *Proceedings of the 52nd Annual Meeting of the Association for Computational Linguistics*, 2014
- Levy, O. & Goldberg, Y.: Neural Word Embedding as Implicit Matrix Factorization, in *Advances in Neural Information Processing Systems 27*, 2014
- Arora, S., Li, Y., Liang, Y., Ma, T. & Risteski, A.: A Latent Variable Model Approach to PMI-based Word Embeddings, in *Transactions of the Association for Computational Linguistics*, 2016
- Stratos, K., Collins, M. & Hsu, D.: Model-based Word Embeddings from Decompositions of Count Matrices, in *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics*, 2015
- Faruqui, M., Tsvetkov, Y., Rastogi, P., Dyer, C.: Problems With Evaluation of Word Embeddings Using Word Similarity Tasks. [CoRRabs/1605.02276](https://arxiv.org/abs/1605.02276), 2016
- Ma, X. & Hovy, E.: End-to-end Sequence Labeling via Bi-directional LSTM-CNNs-CRF, *arXiv:1603.01354*, 2016